

Case Study



BUSINESS TASK


Cyclistic built its successful bike sharing program by offering two methods of customer engagements, casual riders and members. Casual riders are defined as those customers who purchase rides on short term basis, one trip or a day pass. Full members purchase annual memberships. Our financial analyst team believe that we can be more profitable if we increase annual memberships.

In order to increase our memberships, we are tasked with the need to understand our members. We need to understand the differences between annual and casual members. We need to understand how each customer group engages with our service. We need to get an insight into how often each group uses our bikes.

PREPARE

Cyclistic provided 36 trip datasets. (See the datasets at this [link](#).) This data publicly provided by Motivate International Inc. under the licensing protocols available at this [link](#). Motivate maintains and culls the data of personally identifiable information such as gender and credit card information.

For this study, I downloaded 12 months of data from September 2020 through September 2021.



Using this Python code snippet, I unzipped all of the files.














Unzip all files:

```
import requests, zipfile, io, os
print("Libraries Imported")

working_dir = 'C:\\Users\\jkhan\\Documents\\Professional
Development\\Google_Analytics\\Capstone\\Data'
os.chdir(working_dir)

for file in os.listdir(working_dir):
    if zipfile.is_zipfile(file):
        with zipfile.ZipFile(file) as item:
            item.extractall()
```

After unzipping, we see that there are 13 csv files for each month.

Name
 202009-divvy-tripdata.csv
 202010-divvy-tripdata.csv
 202011-divvy-tripdata.csv
 202012-divvy-tripdata.csv
 202101-divvy-tripdata.csv
 202102-divvy-tripdata.csv
 202103-divvy-tripdata.csv
 202104-divvy-tripdata.csv
 202105-divvy-tripdata.csv
 202106-divvy-tripdata.csv
 202107-divvy-tripdata.csv
 202108-divvy-tripdata.csv
 202109-divvy-tripdata.csv

Looking at the csv files in excel, we see that there 13 fields for each data observation.

	A	B	C	D	E	F	G	H	I	J	K	L	M
1	ride_id	rideable_type	started_at	ended_at	start_station_name	start_station_id	end_station_name	end_station_id	start_lat	start_lng	end_lat	end_lng	member_casual

1. ride_id – A unique, alphanumeric identifier for each ride
2. rideable_type – A categorical type docked_bike or electric_bike
3. started_at - The date/time without timestamp where the ride started.
4. ended_at - - The date/time without timestamp where the ride ended.
5. start_station_name - The name of the station where the ride started.
6. start_station_id – Alphanumeric identifier for the starting station
7. end_station_name - - The name of the station where the ride ended.
8. end_station_id - Alphanumeric identifier for the ending station
9. start_lat – Latitude for the starting station
10. start_lng – Longitude for the starting station
11. end_lat – Latitude for the ending station
12. end_lng– Longitude for the ending station
13. member_casual – Categorical variable for member type.

Sorting the data by “started_at” shows the earliest date of the data:

C	D	E
started_at	ended_at	start_station_name
9/1/2020 0:00	9/1/2020 0:37	Michigan Ave & Lake St
9/1/2020 0:00	9/1/2020 0:16	Broadway & Wilson Ave
9/1/2020 0:00	9/1/2020 0:16	Broadway & Wilson Ave
9/1/2020 0:02	9/1/2020 0:15	Shore Dr & 55th St
9/1/2020 0:02	9/1/2020 0:17	Southport Ave & Clybourn Ave
9/1/2020 0:02	9/1/2020 0:17	Southport Ave & Clybourn Ave
9/1/2020 0:02	9/1/2020 0:17	Southport Ave & Clybourn Ave
9/1/2020 0:02	9/1/2020 0:13	Halsted St & Roscoe St
9/1/2020 0:02	9/1/2020 0:17	Southport Ave & Clybourn Ave

Also, there is missing data as well

start_station_name	start_station_id
Damen Ave & Cortland St	
Elizabeth (May) St & Fulton St	
Pine Grove Ave & Irving Park Rd	TA1308000022
Austin Blvd & Lake St	
Ashland Ave & Division St	
Elizabeth (May) St & Fulton St	
Wabash Ave & Roosevelt Rd	TA1305000002
Ashland Ave & Division St	
Clark St & Lincoln Ave	
Green St & Madison St	TA1307000120
State St & 95th St	
Sedgwick St & North Ave	TA1307000038
DuSable Lake Shore Dr & North Blvd	LF-005
Indiana Ave & 103rd St	

PROCESS

For processing, I created a PostgreSQL database and a created a table for the data:

```
-- Table: public.cyclistic_bike_data

-- DROP TABLE public.cyclistic_bike_data;

CREATE TABLE IF NOT EXISTS public.cyclistic_bike_data
(
    ride_id character varying(300) COLLATE pg_catalog."default",
    rideable_type character varying(300) COLLATE pg_catalog."default",
    started_at timestamp without time zone,
    ended_at timestamp without time zone,
    start_station_name character varying(300) COLLATE pg_catalog."default",
    start_station_id character varying(300) COLLATE pg_catalog."default",
    end_station_name character varying(300) COLLATE pg_catalog."default",
    end_station_id character varying(300) COLLATE pg_catalog."default",
    start_lat numeric,
    start_lng numeric,
    end_lat numeric,
    end_lng numeric,
    member_casual character varying(300) COLLATE pg_catalog."default"
)
```

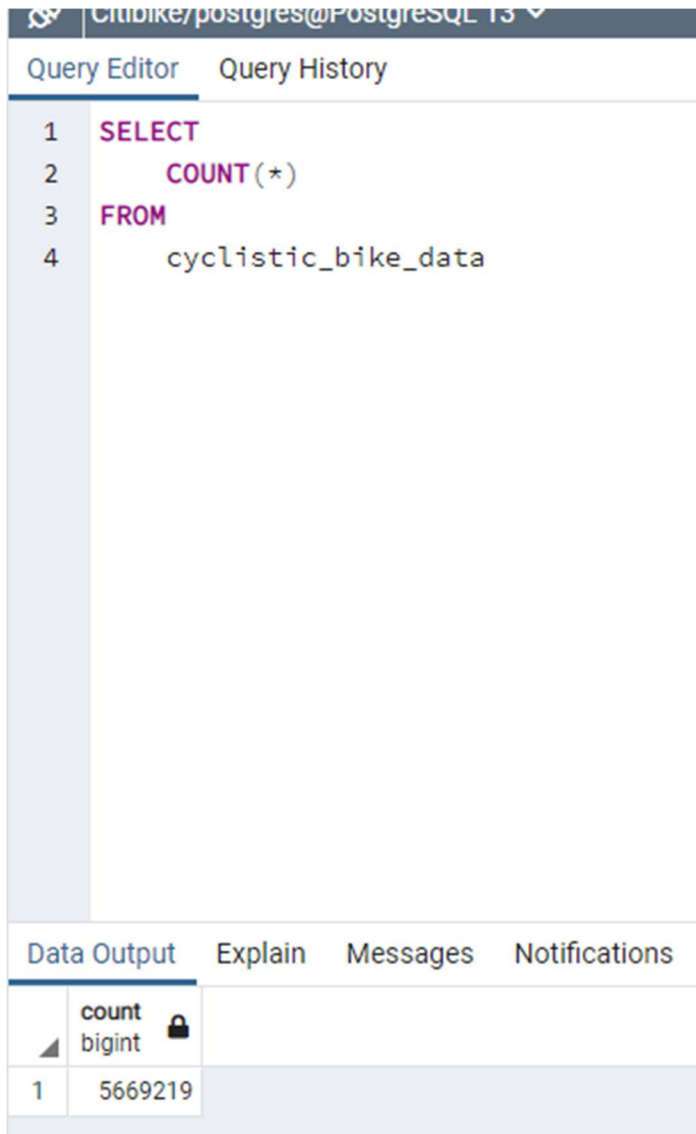
Next, I created a script to load the 13 csv files into the database:

```
psql [psql] 13:
Username [postgres]: postgres
Password for user postgres:
psql (13.4)
WARNING: Console code page (437) differs from Windows code page (1252)
8-bit characters might not work correctly. See psql reference
page "Notes for Windows users" for details.
Type "help" for help.

Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202009-divvy-tripdata.csv' CSV HEADER;
COPY 532958
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202009-divvy-tripdata.csv' CSV HEADER;
COPY 532958
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202010-divvy-tripdata.csv' CSV HEADER;
COPY 388653
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202011-divvy-tripdata.csv' CSV HEADER;
COPY 259716
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202012-divvy-tripdata.csv' CSV HEADER;
COPY 131573
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202101-divvy-tripdata.csv' CSV HEADER;
COPY 96834
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202102-divvy-tripdata.csv' CSV HEADER;
COPY 49622
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202103-divvy-tripdata.csv' CSV HEADER;
COPY 228496
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202104-divvy-tripdata.csv' CSV HEADER;
COPY 337230
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202105-divvy-tripdata.csv' CSV HEADER;
COPY 521633
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202106-divvy-tripdata.csv' CSV HEADER;
COPY 729595
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202107-divvy-tripdata.csv' CSV HEADER;
COPY 822410
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202108-divvy-tripdata.csv' CSV HEADER;
COPY 804352
Citibike=# \copy cyclistic_bike_data FROM 'C:\Users\jkhan\Documents\Professional Development\Google Analytics\Capstone\Data_final\202109-divvy-tripdata.csv' CSV HEADER;
COPY 756147
Citibike=#
```

After import

1. The total number of rides for the past 12 months is **5,669,219**.



The screenshot shows a PostgreSQL Query Editor interface. At the top, the database connection is set to 'Citibike/postgres@PostgreSQL 13'. Below this, there are tabs for 'Query Editor' and 'Query History'. The 'Query Editor' tab is active, displaying a SQL query:

```
1 SELECT
2     COUNT(*)
3 FROM
4     cyclistic_bike_data
```

Below the query editor, there are tabs for 'Data Output', 'Explain', 'Messages', and 'Notifications'. The 'Data Output' tab is active, showing the results of the query in a table format:

	count bigint
1	5669219

2. A count of NULL values shows that there are **814,771** records with at least one null value.

Query Editor		Query History
1	SELECT	
2	COUNT(*)	
3	FROM	
4	cyclistic_bike_data	
5	WHERE	
6	ride_id IS NULL OR	
7	rideable_type IS NULL OR	
8	started_at IS NULL OR	
9	ended_at IS NULL OR	
10	start_station_name IS NULL OR	
11	start_station_id IS NULL OR	
12	end_station_name IS NULL OR	
13	end_station_id IS NULL OR	
14	start_lat IS NULL OR	
15	start_lng IS NULL OR	
16	end_lat IS NULL OR	
17	end_lng IS NULL OR	
18	member_casual IS NULL;	

Data Output		Explain	Messages	Notifications
	count bigint			
1	814771			

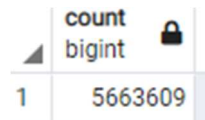
The null values are in features:

start_station_name
start_station_id
end_station_name
end_station_id
start_lat
start_lng
end_lat
end_lng

CLEANING – Part One

KEY DECISION

For the 5,610 rides with no ending station name, id, or geographical information, I removed these 5,610 records from the database leaving a total count of **5,663,609**.



	count
	bigint
1	5663609

For the other null values, I will clean them in R.

3. Adding features to the table:
 - duration – the total time of each ride
 - day_of_week – the day of the week for each ride

[Query Editor](#) [Query History](#)

1
2
3
4
5
6
7
8
9

```
ALTER TABLE public.cyclistic_bike_data
  ADD COLUMN duration time without time zone;

UPDATE public.cyclistic_bike_data
  SET duration = ended_at - started_at;
```

[Data Output](#) [Explain](#) [Messages](#) [Notifications](#)

UPDATE 5663609

Query returned successfully in 1 min 22 secs.

Query Editor		Query History
1	ALTER TABLE public.cyclistic_bike_data	
2	ADD COLUMN day_of_week integer;	
3		
4	UPDATE cyclistic_bike_data	
5	SET day_of_week = DATE_PART('DOW', started_at);	
6		
7	SELECT	
8	day_of_week	
9	FROM	
10	cyclistic_bike_data	
11	LIMIT 50;	
12		
Data Output		Explain Messages Notifications
	day_of_week integer	
31	5	
32	0	
33	5	
34	1	
35	6	
36	6	
37	1	
38	0	
39	2	
40	0	
41	5	
42	1	
43	0	
44	2	

4. Basic Metrics

A. Max, Min, Avg, and Median

Query Editor Query History

```
1 SELECT
2     MAX(duration) AS maximum_duration,
3     AVG(duration) AS average_duration,
4     MIN(duration) AS minimum_duration,
5     PERCENTILE_CONT(0.5) WITHIN GROUP(ORDER BY duration) AS median_duration
6 FROM
7     cyclistic_bike_data;
```

The results from the query above show that the

Maximum duration for a ride is 23:59:59

Minimum duration for a ride is 00:00:00

Average duration is 00:21:52

Median duration is 00:12:44

B. Days of the Week

```
Query Editor  Query History
1  WITH days_map AS(
2      SELECT
3          day_of_week,
4          CASE
5              WHEN day_of_week = 0 THEN 'Sunday'
6              WHEN day_of_week = 1 THEN 'Monday'
7              WHEN day_of_week = 2 THEN 'Tuesday'
8              WHEN day_of_week = 3 THEN 'Wednesday'
9              WHEN day_of_week = 4 THEN 'Thursday'
10             WHEN day_of_week = 5 THEN 'Friday'
11             WHEN day_of_week = 6 THEN 'Saturday'
12         END days
13     FROM
14         cyclistic_bike_data)
15 SELECT
16     days,
17     COUNT(days) AS no_of_rides
18 FROM
19     days_map
20 GROUP BY
21     days
22 ORDER BY no_of_rides DESC
23
```

This query returns the days of the week and the number of rides for each day. Saturdays, Sundays, and Fridays are the most popular days:

	Data Output	Explain	Message
	days text	no_of_rides bigint	
1	Saturday	1012485	
2	Sunday	872585	
3	Friday	821367	
4	Wednesday	763437	
5	Thursday	760171	
6	Tuesday	726128	
7	Monday	707436	

C. Roundtrips

Query Editor Query History

```
1 WITH roundtrips AS (  
2   SELECT  
3     start_station_name,  
4     end_station_name,  
5     SUM(duration) AS time,  
6     COUNT(*) AS roundtrip  
7   FROM  
8     cyclistic_bike_data  
9   WHERE  
10    (start_station_name IS NOT NULL OR  
11    end_station_name IS NOT NULL)  
12    AND  
13    start_lat = end_lat  
14  GROUP BY  
15    start_station_name,  
16    end_station_name)  
17  SELECT  
18    ROW_NUMBER() OVER(ORDER BY roundtrip DESC),  
19    start_station_name,  
20    end_station_name,  
21    roundtrip AS no_of_round_trips_taken,  
22    time/roundtrip AS average_duration  
23  FROM  
24    roundtrips  
25  LIMIT 50;
```

The above query calculates the top 50 most popular round trips where the start and ending stations are the same. The Streeter Drive and Grand Ave round trip is almost twice as popular as the next round trip.

Data Output Explain Messages Notifications

	row_number bigint	start_station_name character varying (300)	end_station_name character varying (300)	no_of_round_trips_taken bigint	average_duration interval
1	1	Streeter Dr & Grand Ave	Streeter Dr & Grand Ave	12067	00:48:38.344659
2	2	Lake Shore Dr & Monroe St	Lake Shore Dr & Monroe St	6683	00:45:38.073021
3	3	Millennium Park	Millennium Park	5851	00:50:16.900701
4	4	Michigan Ave & Oak St	Michigan Ave & Oak St	5740	00:52:28.007317
5	5	Theater on the Lake	Theater on the Lake	4536	00:50:33.329145

D. Single Trips

Query Editor Query History

```
1 WITH trips AS (  
2   SELECT  
3     start_station_name,  
4     end_station_name,  
5     SUM(duration) AS time,  
6     COUNT(*) AS trips_taken  
7   FROM  
8     cyclistic_bike_data  
9   WHERE  
10    start_station_name IS NOT NULL  
11    AND  
12    end_station_name IS NOT NULL  
13    AND  
14    start_station_name != end_station_name  
15  GROUP BY  
16    start_station_name,  
17    end_station_name)  
18  SELECT  
19    ROW_NUMBER() OVER(ORDER BY trips_taken DESC),  
20    start_station_name,  
21    end_station_name,  
22    trips_taken,  
23    time/trips_taken AS avg_trip_length  
24  FROM  
25    trips  
26  
27  LIMIT 50;
```

Data Output						Explain	Messages	Notifications
	row_number bigint		start_station_name character varying (300)	end_station_name character varying (300)	trips_taken bigint		avg_trip_length interval	
1		1	Lake Shore Dr & Monroe St	Streeter Dr & Grand Ave	3685		00:31:43.062958	
2		2	Ellis Ave & 60th St	Ellis Ave & 55th St	3653		00:06:22.254585	
3		3	Streeter Dr & Grand Ave	Millennium Park	3399		00:34:22.014416	
4		4	Ellis Ave & 55th St	Ellis Ave & 60th St	3275		00:07:58.70687	
5		5	Millennium Park	Streeter Dr & Grand Ave	2989		00:44:42.589495	
6		6	Streeter Dr & Grand Ave	Theater on the Lake	2805		00:32:59.403565	

The above query calculates the top 50 most popular single trips where the start and ending stations are not the same. The most popular are Lake Shore Drive & Monroe to Streeter Dr & Grand Ave, Ellis Ave & 60th Street and Ellis Ave & 55th Street.

E. Casual vs. Members

Query Editor		Query History	
1	SELECT		
2	member_casual,		
3	COUNT(*) AS no_of_casual,		
4	CONCAT(ROUND(COUNT(member_casual)/5669219::numeric*100,3), '%') AS percent_of_riders		
5	FROM		
6	cyclistic_bike_data		
7	GROUP BY		
8	member_casual;		

Data Output				Explain	Messages	Notifications
	member_casual character varying (300)		no_of_casual bigint		percent_of_riders text	
1	casual		2585538		45.607%	
2	member		3078071		54.294%	

The results of the above query which counts the number of casual and member riders shows that there are 9% more members than casual riders.

CLEANING – Part Two

KEY DECISION

There are 809,161 records that have NULL values in the station name or id. These records are incomplete and including them in the analysis may lead to incorrect recommendations. There is no way to impute the null values due to mis-matching values in the latitude and longitude features in the dataset.

As a result, the only way to handle these 809,161 will be to delete them from further analysis.

Query Editor Query History

```
1  SELECT
2      start_station_name,
3      end_station_name,
4      start_station_id,
5      end_station_id
6  FROM
7      cyclistic_bike_data
8  WHERE
9      start_station_name IS NULL
10 OR
11 end_station_name IS NULL
12 OR
13 start_station_id IS NULL
14 OR
15 end_station_id IS NULL
16
```

	start_station_name character varying (300)	end_station_name character varying (300)	start_station_id character varying (300)	end_station_id character varying (300)
1	[null]	[null]	[null]	[null]
2	[null]	[null]	[null]	[null]
3	[null]	[null]	[null]	[null]
4	[null]	[null]	[null]	[null]
5	[null]	[null]	[null]	[null]
6	[null]	[null]	[null]	[null]
7	[null]	[null]	[null]	[null]
8	[null]	[null]	[null]	[null]
9	[null]	[null]	[null]	[null]
10	[null]	[null]	[null]	[null]
11	[null]	[null]	[null]	[null]

CLEANING – Part Three

KEY DECISION

There are 5,333 records where the ended_at time is less than the started_at time. These records will have to be removed in order to preserve the integrity of the data:

Data Output
Explain
Messages
Notifications

	ride_id character varying (300)	rideable_type character varying (300)	started_at timestamp without time zone	ended_at timestamp without time zone	start_s charac
1	AA7A3D20CEB995B7	docked_bike	2020-09-05 12:21:25	2020-09-05 12:21:23	Damer
2	FBCAA29D6351CD76	docked_bike	2020-10-07 20:49:35	2020-10-07 20:49:12	Clark S
3	BA783F72B6FE9871	docked_bike	2020-09-05 16:04:32	2020-09-05 16:04:16	Shore I
4	1F7A089EE71EABAF	electric_bike	2020-09-05 16:11:11	2020-09-05 16:11:10	Shedd
5	2127EB98ED316FB2	electric_bike	2020-09-24 12:25:28	2020-09-24 12:25:27	WATSC
6	28E8822BE101E2F7	electric_bike	2020-09-03 12:20:02	2020-09-03 12:20:01	WATSC
7	30754D787851A9D2	electric_bike	2020-09-09 09:22:37	2020-09-09 09:22:36	WATSC

After removing these records, there are 4,849,115 records which will be analyzed.

ANALYZE

(See RMD Notebook File)

Key Findings

- ❖ Casual riders, who are less than half of all riders, ride mostly on the weekends
- ❖ Casual riders ride for less time during the day than members
- ❖ The starting points and routes for the casual rider occur near tourist destinations

Recommendations

- ✓ Create a new class of memberships by offering weekend only passes
- ✓ The benefit to the casual rider will be increased time on the bike without penalty
- ✓ Prices for the weekend passes should be priced at a discount for two day passes
- ✓ Prices for the day and hour passes on the weekend should be raised in order to boost the sales of the weekend passes